

# An introduction to Libra and LibraBFT consensus algorithm

Maurizio Murgia<sup>1</sup>

<sup>1</sup>University of Trento

De Cifris Athesis Seminar — Trento, October 30

# Outline of the talk

- ▶ Brief overview of Libra blockchain.
- ▶ Background:
  - ▶ Byzantine consensus.
  - ▶ Distributed systems.
- ▶ Basic HotStuff.
- ▶ Chained HotStuff and LibraBFT.

# Motivation

Limits of traditional banking:

- ▶ Transferring money internationally is very expensive and slow.
- ▶ ~ 1.8 billions of people do not have access to banks.

Limits of cryptocurrency:

- ▶ Scalability curse.
- ▶ Relatively fast for international transfers, but local payments...
- ▶ Solutions based on PoW waste energy.
- ▶ Expensive.
- ▶ Highly volatile.

## Goals

Infrastructure for moving money (and other assets). Should be:

- ▶ Global.
- ▶ Open, even to people without access to traditional banks.
- ▶ Instant.
- ▶ Low cost.
- ▶ Stable.


## Libra blockchain (1)

- ▶ 1k transactions per second, with latency of 10 seconds.
- ▶ Libra coin is backed by a reserve of assets designed to give it intrinsic value.
- ▶ User can deploy their own resources, accessible through smart contracts.

## Libra blockchain (2)

- ▶ Libra will be permissioned, at least for the first period.  
Transition to the permissioned setting has been planned to occur within five years.
- ▶ Libra uses a Byzantine Fault Tolerant consensus algorithm.
- ▶ Smart contracts: a new programming language, Move, with first class resources.

## On permissioned BFT-based blockchains



Pros:

- ▶ Fast.
- ▶ Scalable.
- ▶ Block finality.

Cons:

- ▶ Decentralised?
- ▶ Truly open?

# The Libra Association (1)





## The Libra Association (2)

From <https://libra.org>:

- ▶ The initial members of the council are the Founding Members and serve as the network's initial validator node.
- ▶ To be such a node, an entity needs to commit \$10 million to the project.
- ▶ Each \$10 million committed entitles one vote in the council, subject to a cap.
- ▶ The Libra Association Council will prevent related entities from presenting themselves as two distinct Founding Members in order to avoid the circumvention of the above measure.

## Calibra Wallet

Calibra will be the official wallet of Libra coins.

- ▶ Integrated with Facebook products:
  - ▶ Facebook.
  - ▶ Whatsapp.
  - ▶ Instagram.
- ▶ Sending money like sending a message or a photo.
- ▶ A bit like WeChat Pay in China.
- ▶ Potential for a great impact in people lives.

## Miscellaneous on Libra

- ▶ Work in progress.
- ▶ Release plan: half 2020.
- ▶ Open source.
- ▶ Testnet available.

## Blockchain? (1)

From the White Paper:

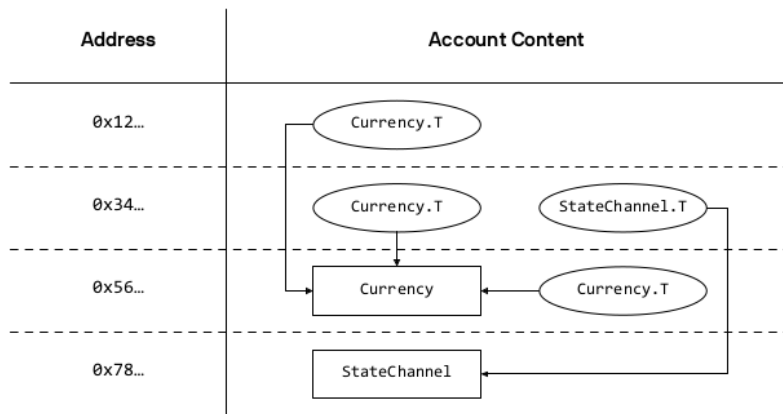
“Unlike previous blockchains, which view the blockchain as a collection of blocks of transactions, the Libra Blockchain is **a single data structure that records the history of transactions and states over time**. This implementation simplifies the work of applications accessing the blockchain, allowing them to read any data from any point in time and verify the integrity of that data using a unified framework.”

## Blockchain? (2)

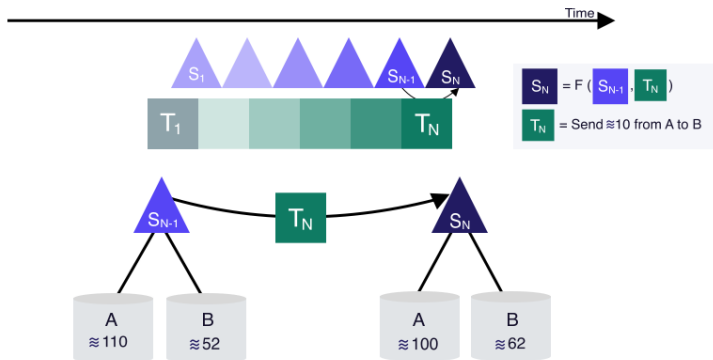
From the report of LibraBFT (State Machine Replication in the LibraBlockchain):

“The leader proposes a new **block consisting of transactions** and sends it to the rest of the validators, who approve the new block if it consists of valid transactions. Once the leader collects a majority of votes, she sends it to the rest of the validators. If a leader fails to propose a valid block or does not aggregate enough votes, a timeout mechanism will force a new round, and a new leader will be chosen from the validators. This way, **new blocks extend the blockchain.**”

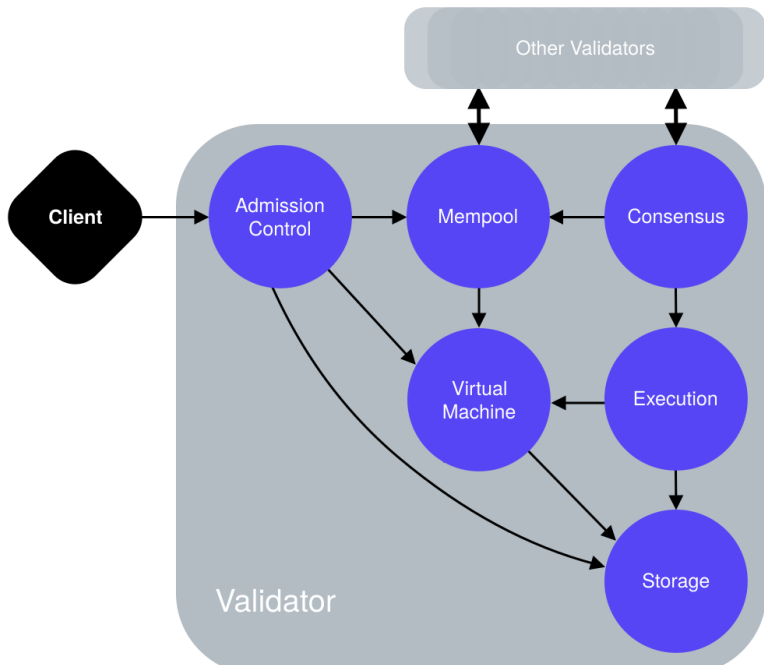
## Blockchain state



# Transactions



## Lifecycle of the Transaction



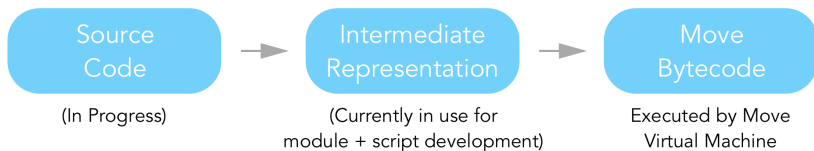


## Smart contracts


- ▶ Libra allows users to create custom resources.
- ▶ Resources might be exchanged, under the rules of the associated smart contract.
- ▶ Resources and smart contracts are implemented in the Move programming language
- ▶ Even Libra coin will be implemented as a smart contract.

# The move language


- ▶ Allows declaration of resources.
- ▶ Ensures that resources are not copied, but just moved.
- ▶ Linear type system.
- ▶ Termination is enforced with a gas mechanism.



# Byzantine Generals Problem (1)



**Coordinated Attack Leading to Victory**



**Uncoordinated Attack Leading to Defeat**

## The Byzantine Generals Problem (2)

- ▶ Participants can communicate only through message passing.
- ▶ Byzantine fault: faulty participants (traitors) can behave arbitrarily, e.g. sending misleading messages.
- ▶ Correctness:
  - ▶ IC1: All loyal lieutenants obey the same order.
  - ▶ IC2: If the commanding general is loyal, then every loyal lieutenant obeys the order he sends.
- ▶ No solution is possible if at least one third of participants are traitors (in partially synchronous networks).

## State Machine Replication (SMR)

According to Wikipedia: State machine replication [...] is a general method for implementing a fault-tolerant service by replicating servers and coordinating client interactions with server replicas.

- ▶ All replicas start with the same initial values.
- ▶ All replicas receiving the same request shall output the same execution result and end up in the same state.
- ▶ Byzantine version: replicas are subject to byzantine faults.
- ▶ The main challenge is guaranteeing that replicas execute the same requests in the same order.

## Network synchrony

Synchronous networks:

- ▶ Operations are coordinated in rounds.
- ▶ Message delivery time is bounded by a fixed delay  $\Delta$ .
- ▶ Not realistic!

Asynchronous networks:

- ▶ Processes can coordinate only with messages.
- ▶ Message delivery is guaranteed, but in unbounded time.
- ▶ Most general setting, consensus problem proved impossible for deterministic algorithms.

Partially synchronous networks:

- ▶ Network delay is unbounded.
- ▶ Long enough runs approximate synchronous runs.
- ▶ Basically, assumes that network faults are eventually fixed and DoS attacks eventually terminates: reasonable in practice!

## The GST model

- ▶ A simple partially synchronous model.
- ▶ Network is asynchronous, until Global Stabilisation Time (GST).
- ▶ After GST, network behaves synchronously.

## Correctness of distributed systems

Properties of distributed systems can be categorised as follows:

- ▶ Safety properties: something (bad) never happens.
- ▶ Liveness properties: something (good) eventually happens.

Safety and liveness in blockchain BFT consensus:

- ▶ Safety: two correct replicas never commit conflicting blocks.
- ▶ Liveness: a new block is eventually produced (after GST).

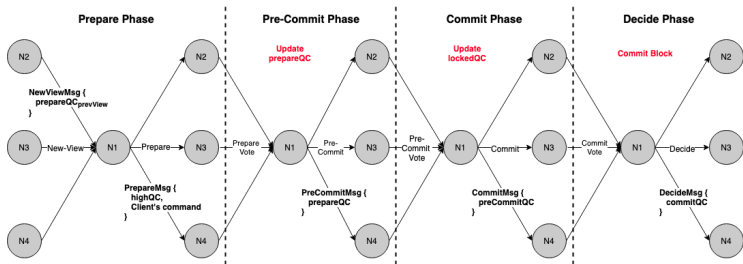


- ▶ Recent BFT consensus algorithm.
- ▶ Fast leader election.
- ▶ Linear communication cost.
- ▶ Optimistic responsiveness: after GST, if the current leader is correct, consensus is reached at the pace of actual (vs. maximum) network delay.

## Preliminaries

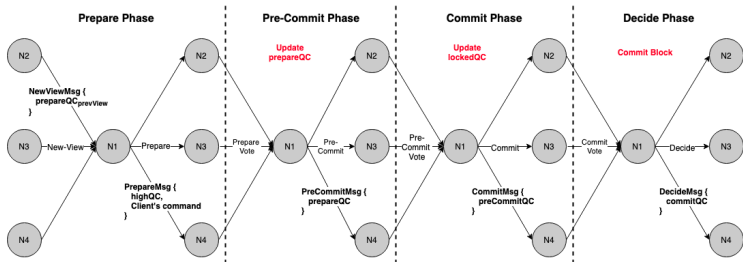
- ▶  $f$  number of byzantine replicas.
- ▶  $n$  total number of replicas.
- ▶ **BFT assumption:**  $n \geq 3f + 1$ .
- ▶ Fact: let  $A$  and  $B$  be two sets of replicas, with at least  $2f + 1$  elements each. Then  $A \cap B$  contains at least one correct replica.
- ▶ Quorum Certificate (QC): a set of  $2f + 1$  (signed) votes of distinct replicas.
- ▶ View: a progressive number; each view has an associated leader.

# Basic HotStuff




- ▶ Basic version of HotStuff, composed of three rounds.
- ▶ Note that replicas only speaks with the current leader.

# Prepare Phase



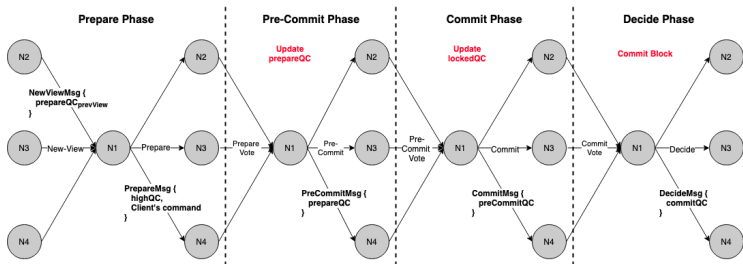
- ▶ **Leader:**
  - ▶ waits for  $2f + 1$  NEW-VIEW messages (sent at the end of previous round), and selects the QC with maximal view number.
  - ▶ creates a new block, with parent the selected block, and broadcasts it in a PREPARE message.
- ▶ **Replica:**
  - ▶ waits for a PREPARE message, and accepts block if:
    - ▶ block extends from lockedQC (safety rule), or
    - ▶  $m.QC.view > lockedQC.view$  (liveness rule).
  - ▶ if message was accepted, sends vote.

# Pre-Commit Phase



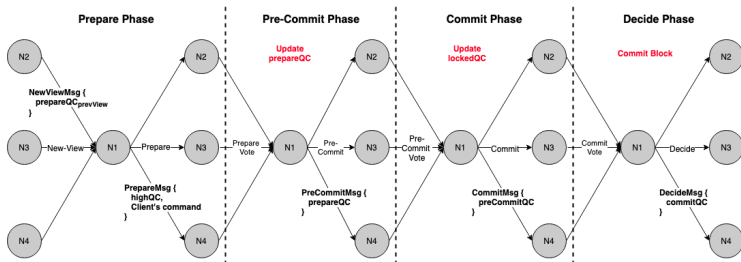
- ▶ **Leader:**
  - ▶ waits for  $2f + 1$  PREPARE votes.
  - ▶ broadcasts a PRE-COMMIT message, with QC of received votes.
- ▶ **Replica:**
  - ▶ waits for a PRE-COMMIT message.
  - ▶ updates prepareQC.
  - ▶ sends vote.

# Commit Phase



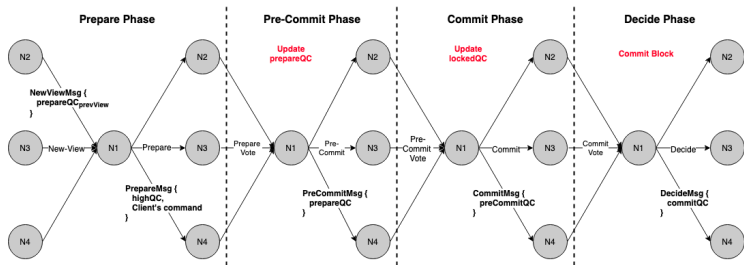
- ▶ **Leader:**
  - ▶ waits for  $2f + 1$  PRE-COMMIT votes.
  - ▶ broadcasts a COMMIT message, with QC of received votes.
- ▶ **Replica:**
  - ▶ waits for a COMMIT message.
  - ▶ updates lockedQC.
  - ▶ sends vote.

# Decide Phase



- ▶ **Leader:**
  - ▶ waits for  $2f + 1$  COMMIT votes.
  - ▶ broadcasts a DECIDE message, with QC of received votes.
- ▶ **Replica:**
  - ▶ waits for a DECIDE message.
  - ▶ increments current view, and sends a NEW-VIEW message.

# Timeouts



- ▶ If an expected message does not arrive, a timeout is triggered.
- ▶ A timedout replica increments current view, and sends a NEW-VIEW message.



## Safety of Basic HotStuff

- ▶ Fact 1: two conflicting QC of the same type must refer to different views.
- ▶ Otherwise, a correct replica voted twice in the same phase: contradiction!

### Theorem (Safety)

*If two blocks are conflicting, they cannot be both committed, each by a correct replica.*

## Safety “proof”

- ▶ By contradiction: suppose  $b1$  and  $b2$  are conflicting and committed. They must have two QC  $qcCommit1$  and  $qcCommit2$  of type COMMIT.
- ▶ By Fact 1 they have different view numbers, say  $v1$  and  $v2$ .
- ▶ Suppose w.l.o.g.  $v1 < v2$ .
- ▶ Let  $qcPrepare3$  be the minimal QC of type PREPARE with view number  $v1 < v3 \leq v2$  and conflicting with  $b1$ .
- ▶ Let  $r$  be a correct replica that voted both in  $qcCommit1$  and  $qcPrepare3$ .
- ▶ At view  $v1$ ,  $r$  updated her lockedQC to a precommitQC on  $b1$ .
- ▶ When  $r$  reached the PREPARE message that she voted in  $qcPrepare3$ , her lockedQC was still a precommitQC on  $b1$  (due to minimality).
- ▶ Then both safety and liveness rules does not hold, and  $r$  could not vote for  $qcPrepare3$ : contradiction!

## On liveness

- ▶ Fact 2: If a correct replica  $i$  locked such that  $\text{lockedQC} = \text{precommitQC}$ , then at least  $f + 1$  correct replicas voted for some  $\text{prepareQC}$  matching  $\text{lockedQC}$ .
- ▶  $\text{PrepareQC}$  was voted by at least  $2f + 1$  replicas. The thesis follows as there are at most  $f$  byzantine replicas.

### Theorem

*After GST, there exists a bounded time period  $T_f$  such that if all correct replicas remain in view  $v$  during  $T_f$  and the leader for view  $v$  is correct, then a decision is reached.*

## “Proof”

- ▶ At the beginning of the round, the leader receives  $2f + 1$  NEW-VIEW, calculates its highQC and sends a PREPARE message.
- ▶ Let  $\text{precommitQC}^*$  be the highest lockedQC among correct replicas.
- ▶ By Fact 2 there are at least  $f + 1$  correct replicas that voted for a  $\text{prepareQC}^*$  matching  $\text{precommitQC}^*$ , and have sent it to the leader in the NEW-VIEW message.
- ▶ The leader must have received (at least) one of such messages, and hence highQC matches  $\text{precommitQC}^*$ .
- ▶ By assumption, all replicas are in view  $v$  and the leader is correct.
- ▶ Then, all correct replicas vote in the PREPARE phase, because liveness rule holds.
- ▶ As we are after GST, all the phases terminate in bounded time, which can be set to be  $T_f$ .

## The pacemaker


The pacemaker of HotStuff is a module which handles leader selection and timeout duration. By the above theorem, in order to obtain liveness, we need that:

- ▶ A correct leader is eventually elected.
- ▶ All correct replicas eventually stay in the same view for  $T_f$ .

They can be achieved as follows:


- ▶ Leader selection must guarantee that every replica eventually becomes a leader: any deterministic mapping that rotates all leaders is fine.
- ▶ At each round, double timeout duration (and start from a non zero value!).

# Chained HotStuff



- ▶ A vote for a block is an implicit vote for parent block.
- ▶ Three chain rule: a block (and its ancestors) is final if it is at the head of three contiguous blocks.

# LibraBFT



- ▶ Based on Chained HotStuff.
- ▶ Several low level details.
- ▶ Some changes in the chain structure.
- ▶ Supports updates of validators.
- ▶ Customised Pacemaker.
- ▶ Leader selection: work in progress, probably based on Verifiable Random Functions.

Thanks!